

# CIFAR-10 Image Classification with ResNet18

MLOps Lab-2, Worksheet-1

Deep Learning Experiment with WandB Integration

**Laksh Mendpara**

Roll Number: B23CS1037

Indian Institute of Technology Jodhpur

February 1, 2026

## Abstract

This report presents a comprehensive deep learning experiment for image classification on the CIFAR-10 dataset using a ResNet18-style Convolutional Neural Network (CNN). The experiment implements a complete MLOps pipeline featuring custom data loaders with advanced augmentation, `torch.compile()` optimization, mixed precision training, and full experiment tracking with Weights & Biases (WandB). We analyze gradient flow dynamics, weight distribution evolution, and training convergence patterns. The model achieves 81.33% test accuracy with 11.17M parameters and 557.89M FLOPs. All visualizations and metrics are logged to WandB for reproducibility and analysis.

## Contents

# 1 Introduction

Deep learning has revolutionized computer vision, with Convolutional Neural Networks (CNNs) achieving remarkable performance on image classification tasks. This experiment implements a complete training pipeline for the CIFAR-10 dataset, emphasizing:

- **Reproducibility:** Full experiment tracking with WandB
- **Visualization:** Gradient flow and weight distribution analysis
- **Optimization:** `torch.compile()` and mixed precision training
- **Best Practices:** Proper data augmentation and regularization

## 1.1 Links

- **Google Colab Notebook:** <https://colab.research.google.com/drive/1ILCjEgsHKy5LvfBMusp=sharing>
- **WandB Dashboard:** <https://wandb.ai/b23cs1037-iit-jodhpur/cifar10-cnn-mlops/runs/6tbudx0x>

# 2 Dataset: CIFAR-10

CIFAR-10 is a widely-used benchmark dataset for image classification consisting of 60,000  $32 \times 32$  color images across 10 classes.

## 2.1 Dataset Statistics

Table 1: CIFAR-10 Dataset Split

Split	Samples
Training Set	45,000
Validation Set	5,000
Test Set	10,000

## 2.2 Classes

The 10 classes are: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, and truck.

## 2.3 Data Augmentation

To improve generalization, we apply the following augmentation pipeline:

Listing 1: Data Augmentation Pipeline

```
transforms.Compose([
    RandomCrop(32, padding=4),
    RandomHorizontalFlip(p=0.5),
```

```

RandomRotation(degrees=15),
ColorJitter(brightness=0.2, contrast=0.2,
            saturation=0.2, hue=0.1),
RandomAffine(degrees=0, translate=(0.1, 0.1)),
ToTensor(),
Normalize(mean=[0.4914, 0.4822, 0.4465],
          std=[0.2470, 0.2435, 0.2616]),
RandomErasing(p=0.25, scale=(0.02, 0.2)),
])

```

## 3 Model Architecture

### 3.1 ResNet18 for CIFAR-10

We use a ResNet18 architecture adapted for CIFAR-10's  $32 \times 32$  input size. The standard ResNet18 (designed for  $224 \times 224$  ImageNet images) is modified as follows:

Table 2: Architecture Modifications for CIFAR-10

Component	ImageNet ResNet18	CIFAR-10 ResNet18
First Conv	$7 \times 7$ , stride 2	$3 \times 3$ , stride 1
Max Pooling	$3 \times 3$ , stride 2	Removed
Input Size	$224 \times 224$	$32 \times 32$

### 3.2 Model Specifications

Table 3: Model Complexity Metrics

Metric	Value
Total Parameters	11,173,962
Trainable Parameters	11,173,962
FLOPs (per image)	557.889M
Architecture	4 stages with [2, 2, 2, 2] blocks

### 3.3 Architecture Details

The ResNet18 architecture consists of:

1. **Initial Convolution:**  $3 \times 3$  conv, 64 filters, BatchNorm, ReLU
2. **Layer 1:** 2 BasicBlocks, 64 filters
3. **Layer 2:** 2 BasicBlocks, 128 filters (stride 2)
4. **Layer 3:** 2 BasicBlocks, 256 filters (stride 2)
5. **Layer 4:** 2 BasicBlocks, 512 filters (stride 2)

6. **Classifier:** AdaptiveAvgPool2d → Linear(512, 10)

Each BasicBlock contains:

- Two  $3 \times 3$  convolutional layers with BatchNorm
- ReLU activation
- Residual (skip) connection

## 4 Training Configuration

### 4.1 Hyperparameters

Table 4: Training Hyperparameters

Hyperparameter	Value
Epochs	25
Batch Size	128
Optimizer	AdamW
Base Learning Rate	0.001
Max Learning Rate	0.01
Weight Decay	$1 \times 10^{-4}$
Gradient Clipping	1.0
LR Scheduler	OneCycleLR

### 4.2 Optimization Techniques

1. **torch.compile()**: PyTorch 2.0+ compilation for optimized execution
2. **Mixed Precision (AMP)**: Automatic Mixed Precision with GradScaler for faster training
3. **Gradient Clipping**: Maximum gradient norm of 1.0 to prevent exploding gradients
4. **OneCycleLR**: Learning rate warm-up and annealing schedule

## 5 Results

### 5.1 Final Performance

Table 5: Final Model Performance

Metric	Value
Best Validation Accuracy	81.60% (Epoch 22)
Final Test Accuracy	81.33%
Final Test Loss	1.0984
Final Training Accuracy	99.07%
Training Time	~28 minutes

### 5.2 Training Progression

Table 6: Training Metrics per Epoch (Selected)

Epoch	Train Loss	Train Acc	Val Loss	Val Acc
1	1.2536	54.20%	0.9954	63.84%
5	0.2801	90.14%	0.7688	75.34%
10	0.0764	97.40%	0.8548	78.54%
15	0.0469	98.38%	0.8927	80.68%
20	0.0382	98.67%	1.0772	79.76%
22	0.0297	98.98%	0.9833	<b>81.60%</b>
25	0.0277	99.07%	1.0877	79.70%

### 5.3 Training Curves




Figure 1: Loss curves over 25 epochs. Training loss decreases smoothly while validation loss shows overfitting after epoch 5.




Figure 2: Accuracy curves over 25 epochs. Training accuracy reaches 99% while validation plateaus at  $\sim 81\%$ .

## 5.4 Sample Predictions




Figure 3: Sample predictions from the test set. Green labels indicate correct predictions (15/16), red indicates misclassification (1/16).

## 6 Gradient Flow Analysis

Gradient flow visualization is crucial for understanding training dynamics and identifying potential issues like vanishing or exploding gradients.

### 6.1 Methodology

For each layer with learnable parameters, we compute:

- **Mean Gradient:**  $\bar{g} = \frac{1}{n} \sum_{i=1}^n |g_i|$
- **Max Gradient:**  $g_{max} = \max(|g_1|, |g_2|, \dots, |g_n|)$

### 6.2 Gradient Flow Comparison




Figure 4: Gradient flow comparison between early and late training. Gradients remain stable in the  $10^{-4}$  to  $10^{-1}$  range throughout.

### 6.3 Observations

1. **Healthy Gradient Flow:** Gradients remain bounded between  $10^{-4}$  and  $10^{-1}$  throughout training, indicating stable optimization.
2. **Residual Connections:** The skip connections in ResNet effectively prevent gradient vanishing, as evidenced by consistent gradient magnitudes across all layers.
3. **BatchNorm Effect:** Batch normalization helps maintain uniform gradient magnitudes, visible in the relatively even distribution across layers.
4. **FC Layer Dominance:** The final fully-connected layer shows the highest gradients ( $\sim 10^{-1}$ ), which is expected as it directly receives the classification loss signal.

## 7 Weight Distribution Analysis

Tracking weight distributions reveals how the network's parameters evolve during training and helps identify potential issues.

## 7.1 Weight Distribution Evolution




Figure 5: Weight distribution evolution from initialization to final trained state.

## 7.2 Observations

1. **Convergence:** Weight distributions become tighter (lower variance) as training progresses, indicating model convergence.
2. **Layer-Specific Patterns:**
  - **conv1:** Narrow distribution centered at 0
  - **layer1-layer4:** Gaussian-like distributions with varying widths
  - **fc:** Wider distribution spanning approximately  $[-0.1, 0.1]$
3. **No Weight Explosion:** All weights remain bounded, confirming effective weight decay regularization.
4. **Kaiming Initialization:** Initial distributions reflect Kaiming normal initialization, which is optimal for ReLU networks.

## 8 Weight Update Analysis

Weight update magnitudes reveal the learning dynamics and the effect of the learning rate scheduler.




Figure 6: Weight update magnitudes per layer across all 25 epochs.

## 8.1 Observations

1. **OneCycleLR Pattern:** The characteristic pattern of OneCycleLR is visible—updates increase during warm-up (epochs 1-12) and decrease during annealing (epochs 13-25).
2. **Layer-Wise Behavior:**
  - Early layers show smaller, more stable updates
  - Later layers exhibit larger, more variable updates
3. **Convergence:** Update magnitudes decrease towards the end of training, indicating approach to a local minimum.
4. **No Layer Collapse:** All layers maintain non-zero updates throughout, confirming healthy learning across the network.

## 9 Key Findings and Discussion

### 9.1 What Worked Well

1. **ResNet Architecture:** Residual connections effectively combat vanishing gradients, as confirmed by gradient flow analysis.
2. **Data Augmentation:** The augmentation pipeline (RandomCrop, ColorJitter, RandomErasing) provided meaningful regularization.
3. **OneCycleLR:** The learning rate schedule with warm-up and annealing led to effective training dynamics.
4. **Mixed Precision:** AMP provided  $\sim$ 1.5-2x speedup without accuracy loss.
5. **`torch.compile()`:** PyTorch compilation provided additional optimization benefits.

## 9.2 Overfitting Analysis

The ~18% gap between training accuracy (99%) and validation accuracy (81%) indicates overfitting:

- **Evidence:** Validation loss increases after epoch 5 while training loss continues decreasing
- **Cause:** The model's 11M parameters may be excessive for CIFAR-10's 50K training samples
- **Mitigation:** Early stopping at epoch 22 (best validation accuracy)

## 9.3 Potential Improvements

1. **Stronger Regularization:** Add dropout, label smoothing, or stochastic depth
2. **Advanced Augmentation:** Implement CutMix, MixUp, or AutoAugment
3. **Smaller Model:** Use a more parameter-efficient architecture
4. **Longer Training:** Train for 100+ epochs with cosine annealing

# 10 Conclusion

This experiment successfully implemented a complete deep learning pipeline for CIFAR-10 classification with comprehensive MLOps practices:

- Achieved **81.33% test accuracy** with ResNet18 (11.17M parameters, 557.89M FLOPs)
- Demonstrated **healthy gradient flow** through gradient visualization
- Tracked **weight distribution evolution** showing proper convergence
- Identified **overfitting patterns** through training curve analysis
- Maintained **full experiment reproducibility** via WandB logging

The comprehensive visualization and logging infrastructure provides valuable insights into model behavior and serves as a template for future deep learning experiments.

## 11 Appendix: All Visualizations

## 11.1 Gradient Flow Progression




Figure 7: Gradient flow visualization across all checkpoints.

## 11.2 Weight Distribution Progression




Figure 8: Weight distribution evolution across all checkpoints.

### 11.3 Weight Updates Progression




Figure 9: Weight update magnitudes across all checkpoints.